



Batch Processing I

(v1.00)

Week 4: September 29, 2026

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at <https://lintool.github.io/cs431-2026f/>

This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International
See <https://creativecommons.org/licenses/by-nc-sa/4.0/> for details



Starting point: you want to process 100 billion records

`map` $(k_1, v_1) \rightarrow \text{List}[(k_2, v_2)]$

What's the processing?
(Doesn't matter, it's an arbitrary computation)

What do you do?

Write the **map**

that's it.


(there is no notion of *before*, *after*, ...)

(runtime, execution engine, Hadoop, Spark, etc.)

What does the *framework* do?

Everything else

(yes, there is *before*, *after*, ... the framework orchestrates)



Yes, the mapper is running on some machine

How does each mapper know what records to process?
(later)

Not very interesting, right?

MapReduce

Programmer specifies two functions:


map $(k_1, v_1) \rightarrow \text{List}[(k_2, v_2)]$

reduce $(k_2, \text{List}[v_2]) \rightarrow \text{List}[(k_3, v_3)]$

All values with the same key are sent to the same reducer

What's the **map** and **reduce**?

(Doesn't matter, it's an arbitrary computation)



This involves
the network

What do you do?

Write the **map** and **reduce**
that's it.

(there is no notion of *before*, *after*, ...)

(runtime, execution engine, Hadoop, Spark, etc.)

What does the *framework* do?

Everything else

(yes, there is *before*, *after*, ... the framework orchestrates)

Word Count: Baseline

```
class Mapper {  
  def map(key: Long, value: String) = {  
    for (word <- tokenize(value)) {  
      emit(word, 1)  
    }  
  }  
}  
  
class Reducer {  
  def reduce(key: String, values: Iterable[Int]) = {  
    for (value <- values) {  
      sum += value  
    }  
    emit(key, sum)  
  }  
}
```

What does local aggregation mean?

MapReduce

Programmer specifies ~~two~~^{four} functions:

map $(k_1, v_1) \rightarrow \text{List}[(k_2, v_2)]$

reduce $(k_2, \text{List}[v_2]) \rightarrow \text{List}[(k_3, v_3)]$

All values with the same key are sent to the same reducer

partition $(k', p) \rightarrow 0 \dots p-1$


Often a simple hash of the key, e.g., $\text{hash}(k') \bmod p$

Divides up key space for parallel reduce operations

combine $(k_2, \text{List}[v_2]) \rightarrow \text{List}[(k_2, v_2)]$

Local aggregation: a reduce that runs locally before data shuffling

Used as an optimization to reduce the amount of data shuffling



This involves
the network

What do you do?

Write the **map, reduce, partition, combine**
that's it.

(there is no notion of *before, after, ...*)

(runtime, execution engine, Hadoop, Spark, etc.)

What does the *framework* do?

Everything else

(yes, there is *before, after, ...* the framework orchestrates)

```

class Mapper {
  def map(key: Long, value: String) = {
    for (word <- tokenize(value)) {
      emit(word, 1)
    }
  }
}

class Combiner {
  def reduce(key: String, values: Iterable[Int]) = {
    for (value <- values) {
      sum += value
    }
    emit(key, sum)
  }
}

class Reducer {
  def reduce(key: String, values: Iterable[Int]) = {
    for (value <- values) {
      sum += value
    }
    emit(key, sum)
  }
}

```

For word count, let's use the reducer as the combiner

Word Count: Mapper Histogram

```
class Mapper {  
  def map(key: Long, value: String) = {  
    val counts = new Map()  
    for (word <- tokenize(value)) {  
      counts(word) += 1  
    }  
  
    for ((k, v) <- counts) {  
      emit(k, v)  
    }  
  }  
}
```

Are combiners still needed?

The group by is guaranteed by the framework
The partitioner is something you write

(these are separate things)

Ideal Scaling: Why not?

Communication is unavoidable

Workers need to share intermediate results...

Which requires communication across machines...

Which requires synchronization...

Which kills performance.

Combiners help here!


Skew creates idle workers

Tasks are never divided perfectly evenly...

And even if they are, processing times can be unpredictable...

Which leads to idle workers.

Partitioners help here!



MapReduce

Programmer specifies ~~two~~^{four} functions:

map $(k_1, v_1) \rightarrow \text{List}[(k_2, v_2)]$

reduce $(k_2, \text{List}[v_2]) \rightarrow \text{List}[(k_3, v_3)]$

All values with the same key are sent to the same reducer

partition $(k', p) \rightarrow 0 \dots p-1$

Often a simple hash of the key, e.g., $\text{hash}(k') \bmod p$

Divides up key space for parallel reduce operations

combine $(k_2, \text{List}[v_2]) \rightarrow \text{List}[(k_2, v_2)]$

Local aggregation: a reduce that runs locally before data shuffling

Used as an optimization to reduce the amount of data shuffling

What do you do?

Write the **map, reduce, partition, combine**
that's it.

(there is no notion of *before, after, ...*)

(runtime, execution engine, Hadoop, Spark, etc.)

What does the *framework* do?

Everything else

(yes, there is *before, after, ...* the framework orchestrates)

MapReduce “Runtime”

Handles scheduling

Assigns workers to map and reduce tasks

Handles “data distribution”

Moves processes to data

Handles coordination

Groups and shuffles intermediate data

Handles errors and faults

Detects worker failures and restarts

Which means...

You have limited control over data and execution flow!

All algorithms must be expressed in m, r, c, p

You don't know:

Where mappers and reducers run

When a mapper or reducer begins or finishes

Which input a particular mapper is processing

Which intermediate key a particular reducer is processing

...

Abstractions

Pros

You don't have to worry about it.
You don't need to know what's going on.



Cons

You can't worry about it (even if you wanted to).
You don't know what's going on (even if you wanted to).



Combiner Design

Combiners and reducers share the same method signature

Sometimes, reducers can serve as combiners

Often, not...

Combiners are optional optimizations

Should not affect algorithm correctness

May be run 0, 1, or multiple times

Reducers are guaranteed to run *exactly* once

Can reducers be used as combiners?

Example: find average of integers associated with the same key

Computing the Mean: Version I

```
class Mapper {  
  def map(key: String, value: Int) = {  
    emit(key, value)  
  }  
}  
  
class Reducer {  
  def reduce(key: String, values: Iterable[Int]) {  
    for (value <- values) {  
      sum += value  
      cnt += 1  
    }  
    emit(key, sum/cnt)  
  }  
}
```

Why can't we use reducer as combiner?

Computing the Mean: Version 2

```
class Mapper {
  def map(key: String, value: Int) =
    context.write(key, value)
}
class Combiner {
  def reduce(key: String, values: Iterable[Int]) = {
    for (value <- values) {
      sum += value
      cnt += 1
    }
    emit(key, (sum, cnt))
  }
}
class Reducer {
  def reduce(key: String, values: Iterable[Pair]) = {
    for ((s, c) <- values) {
      sum += s
      cnt += c
    }
    emit(key, sum/cnt)
  }
}
```

Why doesn't this work?

Computing the Mean: Version 3

```
class Mapper {  
  def map(key: String, value: Int) =  
    context.write(key, (value, 1))  
}  
class Combiner {  
  def reduce(key: String, values: Iterable[Pair]) = {  
    for ((s, c) <- values) {  
      sum += s  
      cnt += c  
    }  
    emit(key, (sum, cnt))  
  }  
}  
class Reducer {  
  def reduce(key: String, values: Iterable[Pair]) = {  
    for ((s, c) <- values) {  
      sum += s  
      cnt += c  
    }  
    emit(key, sum/cnt)  
  }  
}
```

Fixed? Yes!

Another Example

Term co-occurrence matrix for a text collection

$M = N \times N$ matrix ($N =$ vocabulary size)

M_{ij} : number of times i and j co-occur in some context
(for concreteness, let's say context = sentence)

Why?

Large Counting Problems

Term co-occurrence matrix for a text collection
= specific instance of a large counting problem

A large event space (number of terms)

A large number of observations (the collection itself)

Goal: keep track of interesting statistics about the events

Basic approach

Mappers generate partial counts

Reducers aggregate partial counts

How do we aggregate partial counts efficiently?

First Try: “Pairs”

Each mapper takes a sentence:

Generate all co-occurring term pairs

For all pairs, emit (a, b) → count

Reducers sum up counts associated with these pairs

Use combiners!

Pairs: Pseudo-Code

```
class Mapper {  
  def map(key: Long, value: String) = {  
    for (u <- tokenize(value)) {  
      for (v <- neighbors(u)) {  
        emit((u, v), 1)  
      }  
    }  
  }  
}
```

```
class Reducer {  
  def reduce(key: Pair, values: Iterable[Int]) = {  
    for (value <- values) {  
      sum += value  
    }  
    emit(key, sum)  
  }  
}
```

Pairs: Pseudo-Code

One more thing...

```
class Partitioner {  
  def getPartition(key: Pair, value: Int, numTasks: Int): Int = {  
    return key.left % numTasks  
  }  
}
```

“Pairs” Analysis

Advantages

Easy to implement, easy to understand

Disadvantages

Lots of pairs to sort and shuffle around (upper bound?)

Fewer opportunities for combiners to work

Another Try: “Stripes”

Idea: group together pairs into an associative array

(a, b) $\rightarrow$ 1
(a, c) $\rightarrow$ 2
(a, d) $\rightarrow$ 5 a $\rightarrow$ { b: 1, c: 2, d: 5, e: 3, f: 2 }
(a, e) $\rightarrow$ 3
(a, f) $\rightarrow$ 2

Each mapper takes a sentence:

Generate all co-occurring term pairs

For each term, emit a $\rightarrow$ { b: count_b, c: count_c, d: count_d ... }

Reducers perform element-wise sum of associative arrays

$$\begin{array}{r} a \rightarrow \{ b: 1, \quad d: 5, e: 3 \} \\ + \quad a \rightarrow \{ b: 1, c: 2, d: 2, \quad f: 2 \} \\ \hline a \rightarrow \{ b: 2, c: 2, d: 7, e: 3, f: 2 \} \end{array}$$

Key idea: data structure to bring together intermediate results

Stripes: Pseudo-Code

```
class Mapper {  
  def map(key: Long, value: String) = {  
    for (u <- tokenize(value)) {  
      val map = new Map()  
      for (v <- neighbors(u)) {  
        map(v) += 1  
      }  
      emit(u, map)      a → { b: 1, c: 2, d: 5, e: 3, f: 2 }  
    }  
  }  
}
```

```
class Reducer {  
  def reduce(key: String, values: Iterable[Map]) = {  
    val map = new Map()  
    for (value <- values) {  
      map += value      a → { b: 1,      d: 5, e: 3 }  
                        + a → { b: 1, c: 2, d: 2,      f: 2 }  
                        -----  
                        a → { b: 2, c: 2, d: 7, e: 3, f: 2 }  
    }  
    emit(key, map)  
  }  
}
```

“Stripes” Analysis

Advantages

- Far less sorting and shuffling of key-value pairs
- Can make better use of combiners

Disadvantages

- More difficult to implement
- Underlying object more heavyweight
- Overhead associated with data structure manipulations
- Fundamental limitation in terms of size of event space

Stripes >> Pairs?

Tradeoffs

Developer code vs. framework

CPU vs. RAM vs. disk vs. network

Number of key-value pairs: sorting and shuffling data across the network

Size and complexity of each key-value pair: de/serialization overhead

Cost of manipulating data structures

Opportunities for local aggregation

Tradeoffs

Pairs:

- Generates a *lot* more key-value pairs
- Fewer combining opportunities
- More sorting and shuffling
- Simple aggregation at reduce

Stripes:

- Generates fewer key-value pairs (but more complex values)
- More opportunities for combining
- Less sorting and shuffling
- More complex (slower) aggregation at reduce

(At scale, stripes are faster)

Expressing everything in m , r , c , p ... quite limiting.


Can we generalize?

(Yes, but next time...)

Deep(-er) Dive: How does this all work?

Where do we place the data?

Where do we place the compute?



Yes, the mapper is running on some machine

How does each mapper know what records to process?
(*now*)

Immutable Truth #1: At scale, you must distribute work across multiple machines.

Immutable Truth #2: At scale, computing components break all the time.

Where do we place the data?

Trick #1: Partition

Trick #2: Replicate

Answer:
GFS! (circa 2003)

Remember: There are no solutions, only tradeoffs!

GFS: Assumptions

Commodity hardware over “exotic” hardware

Scale “out”, not “up”

High component failure rates

Inexpensive commodity components fail all the time

“Modest” number of huge files

Multi-gigabyte files are common, if not encouraged

Files are write-once, mostly appended to

Logs are a common case

Large streaming reads over random access

Design for high sustained throughput over low latency

GFS: Design Decisions

Files stored as chunks

Fixed size (64MB)

Reliability through replication

Each chunk replicated across 3+ chunkservers

Single master to coordinate access and hold metadata

Simple centralized management

No data caching

Little benefit for streaming reads over large datasets

Simplify the API: not POSIX!

Push many issues onto the client (e.g., data layout)

HDFS = GFS clone (same basic ideas)

From GFS to HDFS

Terminology differences:

GFS master = Hadoop namenode

GFS chunkservers = Hadoop datanodes

Implementation differences:

Different consistency model for file appends


Implementation language

Performance

For the most part, we'll use Hadoop terminology...

HDFS Architecture

(Yes, SPOF!)



Namenode Responsibilities

Managing the file system namespace

Holds file/directory structure, file-to-block mapping,
metadata (ownership, access permissions, etc.)

Coordinating file operations

Directs clients to datanodes for reads and writes
No data is moved through the namenode

Maintaining overall health

Periodic communication with the datanodes
Block re-replication and rebalancing
Garbage collection

Deep(-er) Dive: How does this all work?

Where do we place the data?

Where do we place the compute?

Immutable Truth #1: At scale, you must distribute work across multiple machines.

Immutable Truth #2: At scale, computing components break all the time.


Where do we place the compute?

Trick #1: Partition

Trick #2: Replicate

Remember: There are no solutions, only tradeoffs!

Aggregated vs. Disaggregated



Hadoop assumes an aggregated architecture

Aggregated vs. Disaggregated

The case *for* aggregated

Every processor can read from local disk

Less reliance on high-performance interconnects

Basic Cluster Components

Namenode (NN)

Holds HDFS metadata

Jobtracker (JT)


Coordinator for MapReduce jobs

On *each* of the worker machines:

Tasktracker (TT): contains multiple task slots


Datanode (DN): serves HDFS data blocks

Putting everything together...



Key idea: align map tasks with HDFS blocks


Hadoop assumes an aggregated architecture



Yes, the mapper is running on some machine

How does each mapper know what records to process?
(*now*)

Aggregated vs. Disaggregated



Hadoop assumes an aggregated architecture

Aggregated vs. Disaggregated

The case *against* disaggregated

Requires higher-performance interconnects

The case *for* disaggregated


Compute and storage demands grow differently

Compute demands fluctuate; data persists

Heterogeneous data processing engines + storage

Note: even for disaggregated architectures, partitioning is necessary!


Storage Hierarchy




Disaggregation within a Rack

Rack with switch

Rack sharing switch



Reality Today: Hybrid Designs



Examples: Amazon EC2, Google Compute Engine, Azure Virtual Machines

Examples: Amazon S3, Google Cloud Storage, Azure Blob Storage

